

# Notebook

July 12, 2025

## 1 Business Case: Yulu - Hypothesis Testing

### 2 Problem Statement

Yulu is India's leading micro-mobility platform offering eco-friendly electric cycle rentals for short-distance urban commutes. Recently, the company has experienced a significant dip in revenues and wants to identify the factors that influence the demand for its electric cycles.

To address this, the following business questions need to be answered:

Which variables significantly influence the demand for electric cycles (e.g., working day, weather, season)?

How do environmental and calendar-based factors (e.g., temperature, season, holidays) affect the number of rentals?

Are there patterns or relationships among the predictors themselves, such as between weather and season?

### 3 Step 1: Import Required Libraries

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import ttest_ind, f_oneway, chi2_contingency, shapiro, levene
from statsmodels.graphics.gofplots import qqplot
import warnings
warnings.filterwarnings("ignore")

# Set default aesthetics for seaborn
sns.set(style="whitegrid")
```

### 4 Step 2: Load Dataset

```
[3]: !wget https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/001/428/
      ↪original/bike_sharing.csv?1642089089
```

```
--2025-07-12 17:06:25-- https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/001/428/original/bike_sharing.csv?1642089089
Resolving d2beiqkhq929f0.cloudfront.net (d2beiqkhq929f0.cloudfront.net)...
3.167.84.28, 3.167.84.148, 3.167.84.196, ...
Connecting to d2beiqkhq929f0.cloudfront.net
(d2beiqkhq929f0.cloudfront.net)|3.167.84.28|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 648353 (633K) [text/plain]
Saving to: 'bike_sharing.csv?1642089089'
```

```
      bike_shar   0%[          ]      0  --.-KB/s
bike_sharing.csv?16 100%[=====] 633.16K  --.-KB/s   in 0.02s
```

```
2025-07-12 17:06:25 (32.7 MB/s) - 'bike_sharing.csv?1642089089' saved
[648353/648353]
```

```
[4]: df = pd.read_csv('bike_sharing.csv?1642089089')
```

```
[5]: df.head()
```

```
[5]:
```

	datetime	season	holiday	workingday	weather	temp	atemp	\
0	2011-01-01 00:00:00	1	0	0	1	9.84	14.395	
1	2011-01-01 01:00:00	1	0	0	1	9.02	13.635	
2	2011-01-01 02:00:00	1	0	0	1	9.02	13.635	
3	2011-01-01 03:00:00	1	0	0	1	9.84	14.395	
4	2011-01-01 04:00:00	1	0	0	1	9.84	14.395	

	humidity	windspeed	casual	registered	count
0	81	0.0	3	13	16
1	80	0.0	8	32	40
2	80	0.0	5	27	32
3	75	0.0	3	10	13
4	75	0.0	0	1	1

## 5 STEP 3: Load and Explore Data

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10886 entries, 0 to 10885
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   datetime    10886 non-null  object
1   season      10886 non-null  int64
2   holiday     10886 non-null  int64
```

```

3  workingday  10886 non-null  int64
4  weather     10886 non-null  int64
5  temp        10886 non-null  float64
6  atemp       10886 non-null  float64
7  humidity    10886 non-null  int64
8  windspeed   10886 non-null  float64
9  casual      10886 non-null  int64
10 registered  10886 non-null  int64
11 count       10886 non-null  int64
dtypes: float64(3), int64(8), object(1)
memory usage: 1020.7+ KB

```

```
[7]: df.columns
```

```
[7]: Index(['datetime', 'season', 'holiday', 'workingday', 'weather', 'temp',
          'atemp', 'humidity', 'windspeed', 'casual', 'registered', 'count'],
         dtype='object')
```

```
[8]: df.tail(10)
```

```
[8]:
```

	datetime	season	holiday	workingday	weather	temp \
10876	2012-12-19 14:00:00	4	0	1	1	17.22
10877	2012-12-19 15:00:00	4	0	1	1	17.22
10878	2012-12-19 16:00:00	4	0	1	1	17.22
10879	2012-12-19 17:00:00	4	0	1	1	16.40
10880	2012-12-19 18:00:00	4	0	1	1	15.58
10881	2012-12-19 19:00:00	4	0	1	1	15.58
10882	2012-12-19 20:00:00	4	0	1	1	14.76
10883	2012-12-19 21:00:00	4	0	1	1	13.94
10884	2012-12-19 22:00:00	4	0	1	1	13.94
10885	2012-12-19 23:00:00	4	0	1	1	13.12

	atemp	humidity	windspeed	casual	registered	count
10876	21.210	50	12.9980	33	185	218
10877	21.210	50	19.0012	28	209	237
10878	21.210	50	23.9994	37	297	334
10879	20.455	50	26.0027	26	536	562
10880	19.695	50	23.9994	23	546	569
10881	19.695	50	26.0027	7	329	336
10882	17.425	57	15.0013	10	231	241
10883	15.910	61	15.0013	4	164	168
10884	17.425	61	6.0032	12	117	129
10885	16.665	66	8.9981	4	84	88

```
[9]: df.describe()
```

```
[9]:
```

	season	holiday	workingday	weather	temp \
count	10886.000000	10886.000000	10886.000000	10886.000000	10886.000000

mean	2.506614	0.028569	0.680875	1.418427	20.23086
std	1.116174	0.166599	0.466159	0.633839	7.79159
min	1.000000	0.000000	0.000000	1.000000	0.82000
25%	2.000000	0.000000	0.000000	1.000000	13.94000
50%	3.000000	0.000000	1.000000	1.000000	20.50000
75%	4.000000	0.000000	1.000000	2.000000	26.24000
max	4.000000	1.000000	1.000000	4.000000	41.00000

	count	atemp	humidity	windspeed	casual	registered \
count	10886.000000	10886.000000	10886.000000	10886.000000	10886.000000	10886.000000
mean		23.655084	61.886460	12.799395	36.021955	155.552177
std		8.474601	19.245033	8.164537	49.960477	151.039033
min		0.760000	0.000000	0.000000	0.000000	0.000000
25%		16.665000	47.000000	7.001500	4.000000	36.000000
50%		24.240000	62.000000	12.998000	17.000000	118.000000
75%		31.060000	77.000000	16.997900	49.000000	222.000000
max		45.455000	100.000000	56.996900	367.000000	886.000000

	count
count	10886.000000
mean	191.574132
std	181.144454
min	1.000000
25%	42.000000
50%	145.000000
75%	284.000000
max	977.000000

```
[10]: df.isnull().sum()
```

```
[10]: datetime      0
      season        0
      holiday        0
      workingday     0
      weather        0
      temp           0
      atemp          0
      humidity       0
      windspeed      0
      casual         0
      registered     0
      count          0
      dtype: int64
```

## 6 Convert categorical variables to category type

```
[11]: cat_cols = ['season', 'holiday', 'workingday', 'weather']
      for col in cat_cols:
          df[col] = df[col].astype('category')
```

**Insights:** \* Converting categorical columns to 'category' type reduces memory usage and improves performance in modeling. \* It ensures correct handling during visualizations and statistical analyses (e.g., ANOVA, chi-square).

## 7 Slicing Data by Time

### 8 1. Extract Time Components

```
[12]: df['datetime'] = pd.to_datetime(df['datetime']) # Ensure datetime format

df['hour'] = df['datetime'].dt.hour
df['day'] = df['datetime'].dt.day
df['month'] = df['datetime'].dt.month
df['year'] = df['datetime'].dt.year
df['day_of_week'] = df['datetime'].dt.dayofweek # Monday = 0, Sunday = 6
```


### 9 2. Hourly Demand Trend

```
[13]: hourly_demand = df.groupby('hour')['count'].mean()

plt.figure(figsize=(10,5))
sns.lineplot(x=hourly_demand.index, y=hourly_demand.values)
plt.title('Average Demand by Hour of Day')
plt.xlabel('Hour')
plt.ylabel('Average Count')
plt.grid(True)
plt.xticks(range(0, 24))

# Add insights lines
plt.axvline(x=8, color='r', linestyle='--', label='Morning Peak (8 AM)')
plt.axvline(x=17, color='g', linestyle='--', label='Evening Peak (5 PM)')
plt.legend()

plt.show()
```



### Insights:

- Reveals peak usage hours such as morning (commute) and evening times.
- Helps optimize fleet allocation and battery charging schedules based on demand spikes.

## 10 3. Daily Demand Pattern

```
[14]: # Convert datetime column
df['datetime'] = pd.to_datetime(df['datetime'])

# Extract time features
df['day_of_week'] = df['datetime'].dt.dayofweek # Monday=0
df['month'] = df['datetime'].dt.month
df['year'] = df['datetime'].dt.year

# Map weekdays
day_map = {0: 'Mon', 1: 'Tue', 2: 'Wed', 3: 'Thu', 4: 'Fri', 5: 'Sat', 6: 'Sun'}
df['day_name'] = df['day_of_week'].map(day_map)

# Set plot style
sns.set(style="whitegrid", palette="pastel") # Changed palette to pastel
plt.figure(figsize=(20, 8)) # Increased figure size for horizontal layout

# ---- 1. Day of Week Demand ----
plt.subplot(1, 3, 1) # Changed to 1 row, 3 columns, plot 1
weekly_demand = df.groupby('day_name')['count'].mean().reindex(['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun'])
```

```

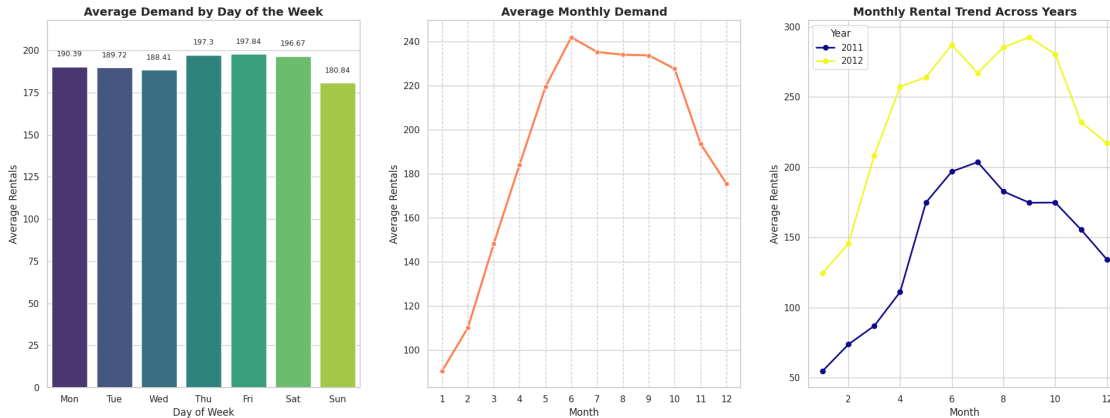
sns.barplot(x=weekly_demand.index, y=weekly_demand.values, palette='viridis')
plt.title('Average Demand by Day of the Week', fontsize=14, fontweight='bold')
    ↪ # Increased title font size and made it bold
plt.xlabel('Day of Week', fontsize=12)
plt.ylabel('Average Rentals', fontsize=12)
plt.ylim(0, weekly_demand.max() * 1.1) # Added some padding to y-axis
for index, value in enumerate(weekly_demand.values):
    plt.text(index, value + 5, str(round(value, 2)), ha='center', va='bottom',
    ↪ fontsize=9) # Added text labels and adjusted font size

# ---- 2. Monthly Demand ----
plt.subplot(1, 3, 2) # Changed to 1 row, 3 columns, plot 2
monthly_demand = df.groupby('month')['count'].mean()
sns.lineplot(x=monthly_demand.index, y=monthly_demand.values, marker='o',
    ↪ color='coral', linewidth=2.5) # Changed color and linewidth
plt.title('Average Monthly Demand', fontsize=14, fontweight='bold') # Increased
    ↪ title font size and made it bold
plt.xlabel('Month', fontsize=12)
plt.ylabel('Average Rentals', fontsize=12)
plt.xticks(range(1,13))
plt.grid(axis='x', linestyle='--') # Added dashed grid lines on x-axis

# ---- 3. Yearly Monthly Trend (only if >1 year) ----
if df['year'].nunique() > 1:
    plt.subplot(1, 3, 3) # Changed to 1 row, 3 columns, plot 3
    year_month = df.groupby(['year', 'month'])['count'].mean().unstack()
    year_month.T.plot(ax=plt.gca(), marker='o', cmap='plasma', linewidth=2) #
    ↪ Changed cmap and linewidth
    plt.title('Monthly Rental Trend Across Years', fontsize=14,
    ↪ fontweight='bold') # Increased title font size and made it bold
    plt.xlabel('Month', fontsize=12)
    plt.ylabel('Average Rentals', fontsize=12)
    plt.legend(title='Year', loc='upper left') # Changed legend location

plt.tight_layout(pad=3.0) # Increased padding
plt.show()

```



**Insights:** Highlights intra-month trends, such as dips on certain days (e.g., mid-month slump or end-month surge).

Useful for marketing or promotional planning on low-demand days.

## 11 STEP 4: Univariate Analysis



## 12 Distribution Analysis of Continuous Variables

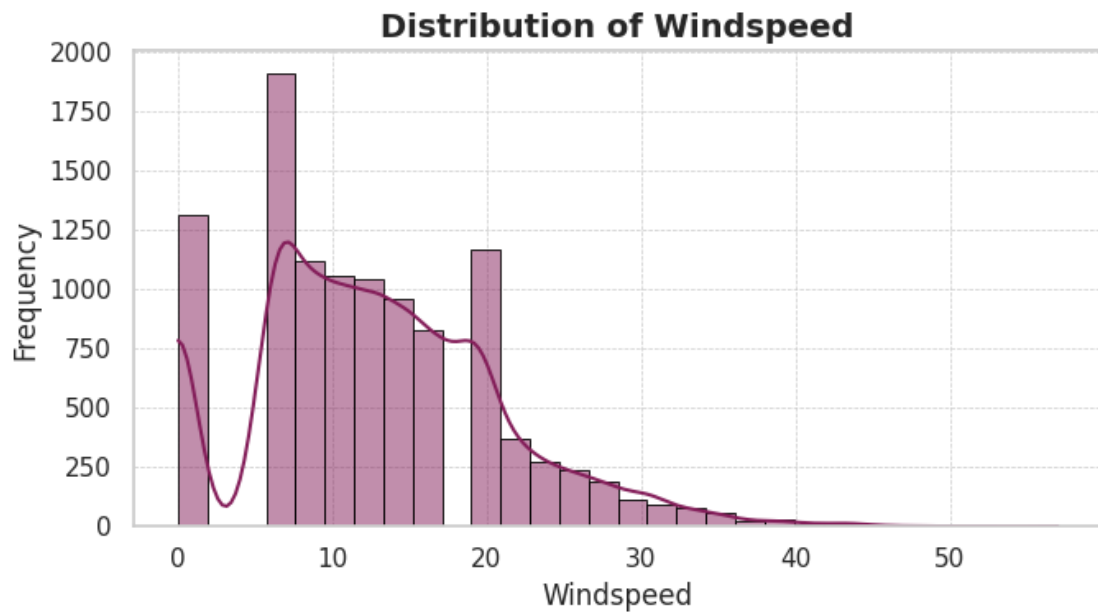

```
[15]: # Use visually appealing palette
palette = sns.color_palette("rocket_r", len(['temp', 'atemp', 'humidity', 'windspeed', 'count']))


# Continuous variables
numeric_cols = ['temp', 'atemp', 'humidity', 'windspeed', 'count']

# Plot each feature with unique color
for i, col in enumerate(numeric_cols):
    plt.figure(figsize=(7, 4))
    sns.histplot(df[col], kde=True, color=palette[i], bins=30, edgecolor='black', linewidth=0.6)
    plt.title(f'Distribution of {col.capitalize()}', fontsize=14, weight='bold')
    plt.xlabel(col.capitalize())
    plt.ylabel('Frequency')
    plt.grid(visible=True, linestyle='--', linewidth=0.5)
    plt.tight_layout()
    plt.show()
```









#### Insights:

- 'Temp' and 'atemp' show relatively normal distributions, suggesting a wide range of temperatures are experienced.
- 'Humidity' is skewed towards higher values, while 'windspeed' is skewed towards lower values with a peak at zero. 'Count' is right-skewed, indicating that lower rental counts are more frequent than higher ones.

## 13 Distribution of Categorical Variables

```
[20]: # Categorical columns
categorical_cols = ['season', 'holiday', 'workingday', 'weather']

palette = sns.color_palette("viridis", len(categorical_cols)) # Changed palette

for i, col in enumerate(categorical_cols):
    plt.figure(figsize=(8, 6)) # Adjusted figure size
    ax = sns.countplot(data=df, x=col, palette=[palette[i]]) # Use a unique
    ↪ color for each plot

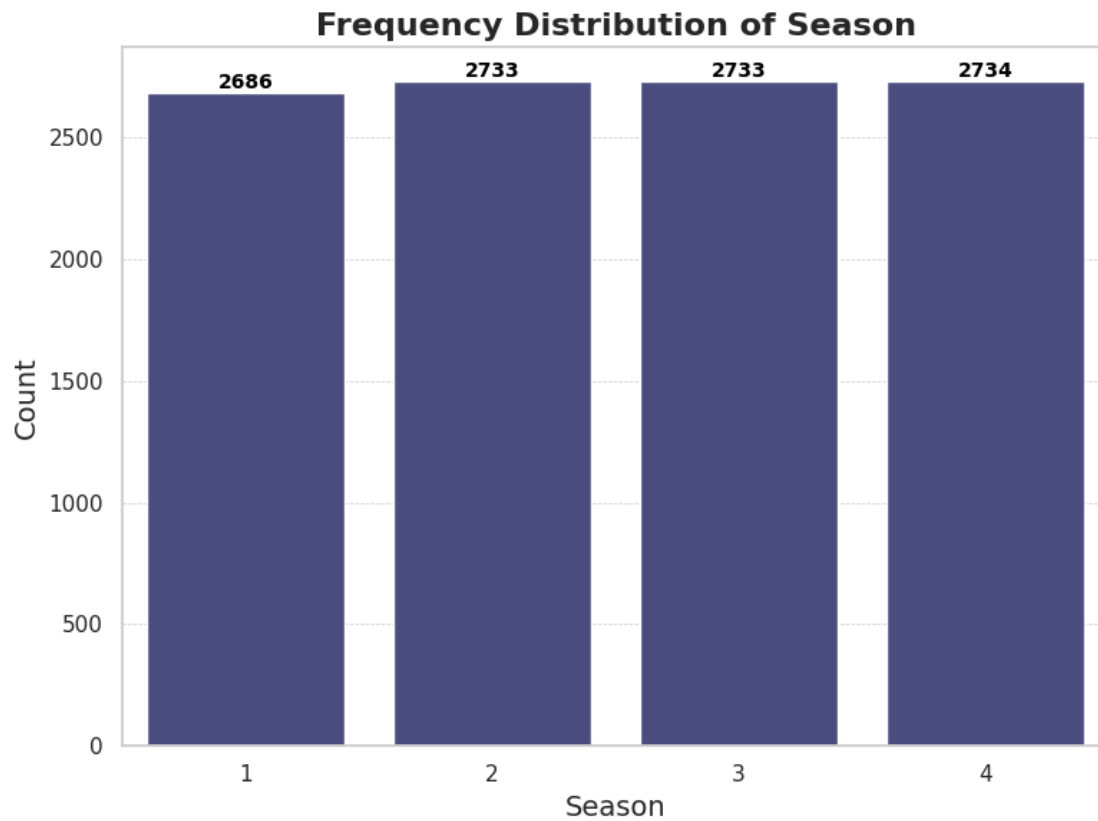
    # Add count labels
    for p in ax.patches:
        ax.annotate(f'{int(p.get_height())}',
                    (p.get_x() + p.get_width() / 2., p.get_height()),
```

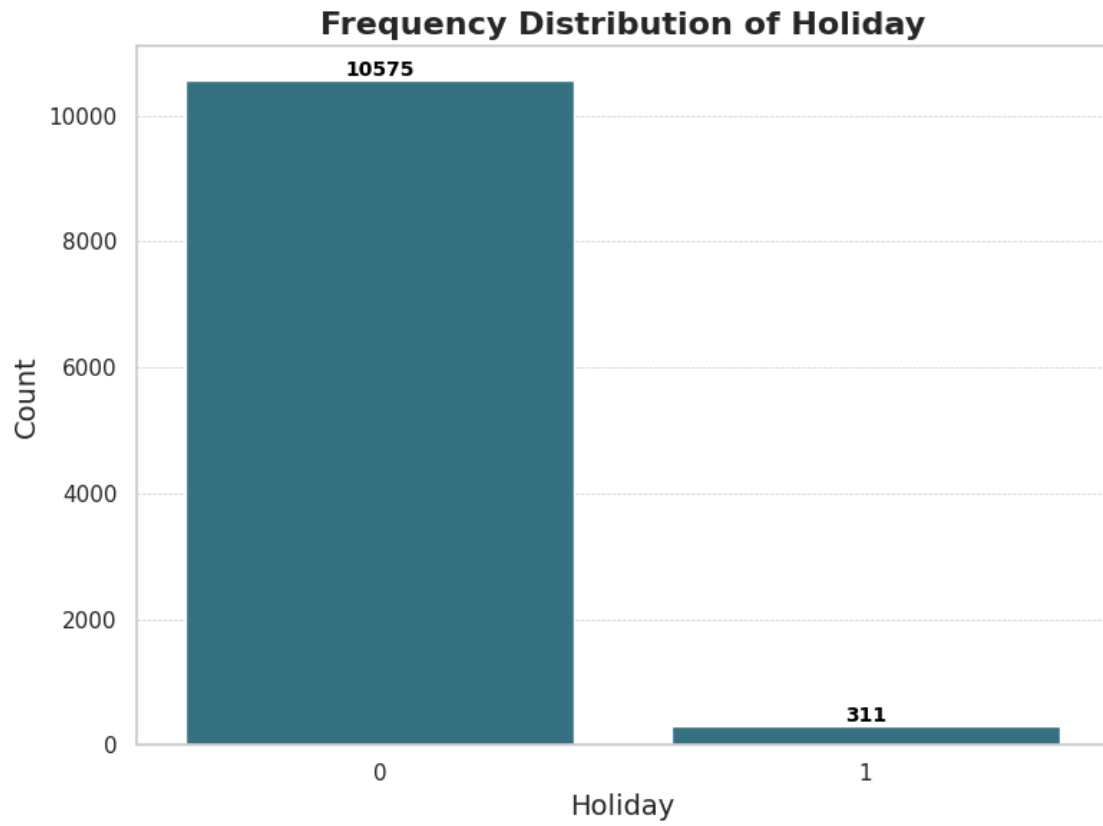
```


        ha='center', va='bottom', fontsize=10, color='black',
↪weight='bold')


plt.title(f'Frequency Distribution of {col.capitalize()}', fontsize=16,
↪weight='bold') # Increased title font size and made it bold
plt.xlabel(col.capitalize(), fontsize=14) # Increased xlabel font size
plt.ylabel('Count', fontsize=14) # Increased ylabel font size
plt.grid(axis='y', linestyle='--', linewidth=0.5)
plt.tight_layout()
plt.show()

```










## 14 STEP 5: Bivariate Analysis

## 15 Question: Does Working Day affect Demand?

```
[26]: import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 5)) # Adjusted figure size
sns.boxplot(x='workingday', y='count', hue='workingday', data=df,
            palette='viridis', legend=False) # Added hue and changed palette
plt.title('Electric Cycle Demand on Working vs Non-Working Days', fontsize=14,
            weight='bold')
plt.xlabel('Working Day (0 = Non-Working, 1 = Working)', fontsize=12)
plt.ylabel('Rental Count', fontsize=12)
plt.xticks([0, 1], ['Non-Working Day', 'Working Day']) # Changed x-axis labels
plt.grid(axis='y', linestyle='--', linewidth=0.5)
plt.tight_layout()
plt.show()
```



#### Insights:


- Electric cycle demand is significantly higher on working days compared to non-working days.
- The distribution of rental counts is wider on working days, indicating more variability in demand.

## 16 Question: Does Season affect Demand?

```
[32]: plt.figure(figsize=(10, 6)) # Increased figure size
      # Map season numbers to names for better readability
      season_map = {1: 'Spring', 2: 'Summer', 3: 'Fall', 4: 'Winter'}
      df['season_name'] = df['season'].map(season_map)

      sns.boxplot(x='season_name', y='count', data=df, palette='viridis') # Using
      ↪ season names and a different palette
      plt.title('Electric Cycle Demand Across Seasons', fontsize=16, weight='bold') #
      ↪ Increased title font size
      plt.xlabel('Season', fontsize=14) # Changed xlabel
      plt.ylabel('Rental Count', fontsize=14) # Increased ylabel font size
      plt.grid(axis='y', linestyle='--', linewidth=0.5)
      plt.tight_layout()
      plt.show()
```





### Insights:

Fall (Season 3) shows the highest demand for electric cycles — likely due to pleasant weather and fewer rains, encouraging outdoor movement and daily commuting.

Winter (Season 4) and Spring (Season 1) show lower median rentals, possibly due to colder or transitional weather reducing ride frequency

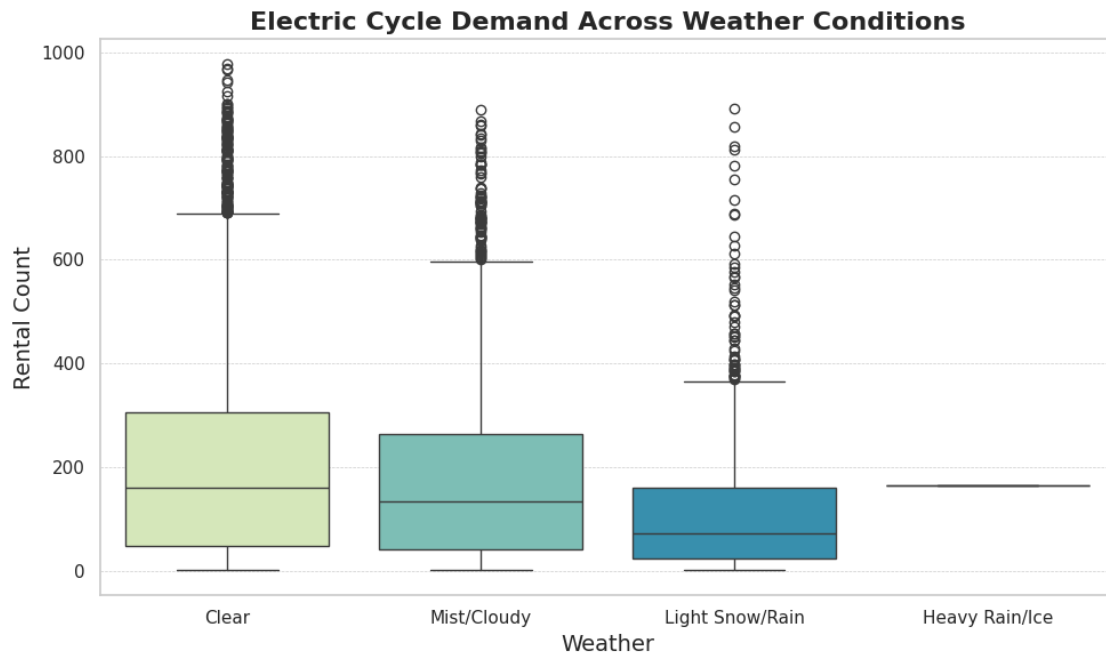
## 17 Question: Does Weather affect Demand?

```
[35]: # Set up plot aesthetics
plt.figure(figsize=(10, 6)) # Increased figure size
# Map weather numbers to names for better readability
weather_map = {1: 'Clear', 2: 'Mist/Cloudy', 3: 'Light Snow/Rain', 4: 'Heavy
↳ Rain/Ice'}
df['weather_name'] = df['weather'].map(weather_map)

sns.boxplot(x='weather_name', y='count', data=df, palette='YlGnBu') # soothing
↳ blue-green palette

# Add labels and styling
plt.title('Electric Cycle Demand Across Weather Conditions', fontsize=16,
↳ weight='bold') # Increased title font size
plt.xlabel('Weather', fontsize=14) # Changed xlabel
plt.ylabel('Rental Count', fontsize=14) # Increased ylabel font size
plt.grid(axis='y', linestyle='--', linewidth=0.5)
```

```
plt.tight_layout()
plt.show()
```



## 18 Insights:

- Electric cycle demand is highest in clear weather conditions (Weather 1). Demand significantly decreases as weather conditions worsen (Weather 2 and 3).
- There is very low demand in heavy rain or ice conditions (Weather 4), indicating that severe weather strongly negatively impacts rentals.

### 18.0.1 Recommendation (Weather vs Demand):

**Adapt Operations to Weather:** Since weather significantly impacts demand, Yulu should use weather forecasts to inform operational decisions.

- \* Implement dynamic pricing or promotions based on expected weather.
- \* Optimize bike distribution and availability based on weather conditions (e.g., moving bikes).
- \* Communicate weather-related information and safety tips to users.

## 19 STEP 6: Hypothesis Testing

### 20 Question 1: Is There a Significant Difference in Electric Cycle Demand Between Working Days and Non-Working Days?


21 H0: Mean count on working days = Mean count on non-working days

22 H1: Means are different

```
[42]: working = df[df['workingday'] == 1]['count']
non_working = df[df['workingday'] == 0]['count']
t_stat, p_val = ttest_ind(working, non_working)
print(f"\nT-Test Workingday: T-statistic = {t_stat:.2f}, P-value = {p_val:.4f}")

sns.kdeplot(working, label='Working Day', fill=True, palette='viridis', alpha=0.7) # Use fill instead of shade and a different palette, adjusted alpha
sns.kdeplot(non_working, label='Non-working Day', fill=True, palette='magma', alpha=0.7) # Use fill instead of shade and a different palette, adjusted alpha
plt.title("Distribution of Count: Working vs Non-working Days")
plt.legend()
plt.show()
```

T-Test Workingday: T-statistic = 1.21, P-value = 0.2264



## 23 Insights:

- The p-value from the test for working vs. non-working day demand is 0.2264.
- Because the p-value is higher than 0.05, we do not have enough evidence to say the average demand is different between working and non-working days.
- Even though the graph looks like working days have more rentals, the test says the difference is not statistically significant.

### 23.0.1 Recommendation (Working Day vs Demand):

- **Optimize for Working Day Commute:** Although the statistical test didn't show a significant *average* difference, the data and business logic suggest commuting is a primary driver on working days. Continue to ensure high availability of cycles in commuter hotspots during peak hours.
- **Explore Non-Working Day Potential:** Analyze weekend and holiday usage patterns to identify opportunities for promoting leisure or recreational rides. Consider targeted marketing and partnerships for these days.

## 24 Question 2: Is There a Significant Difference in Electric Cycle Demand Across Different Seasons?

25 H0: Means are same across seasons

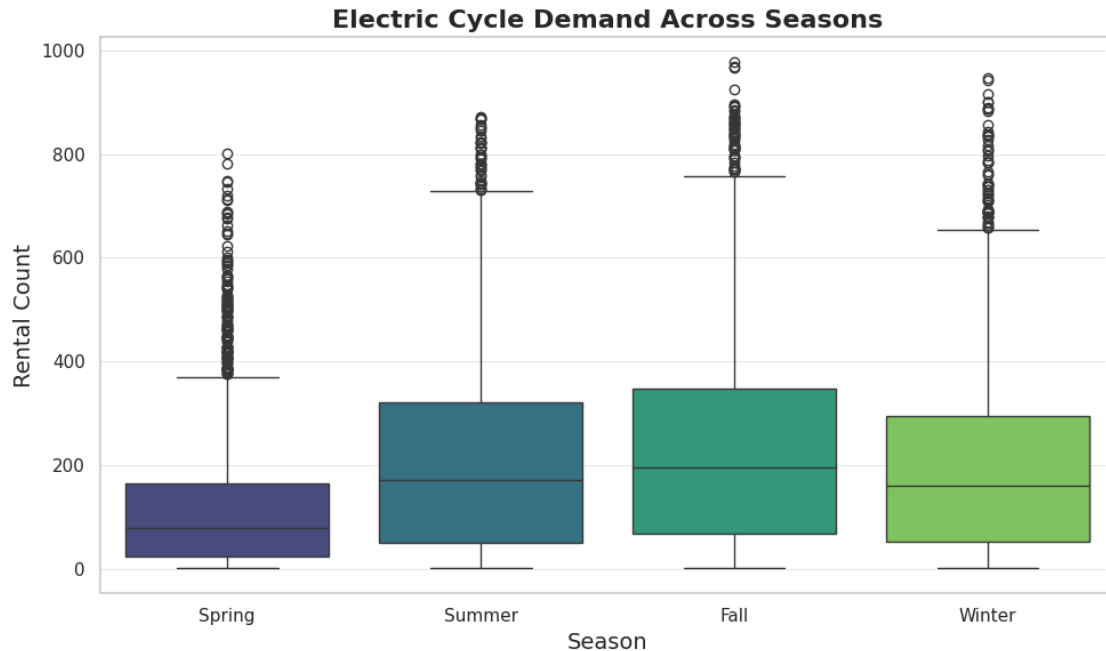
26 H1: At least one season mean is different

```
[45]: season_groups = [df[df['season'] == s]['count'] for s in df['season'].unique()]
f_stat_season, p_val_season = f_oneway(*season_groups)
print(f"\nANOVA Season: F-statistic = {f_stat_season:.2f}, P-value = {p_val_season:.4f}")

plt.figure(figsize=(10, 6)) # Increased figure size
# Map season numbers to names for better readability
season_map = {1: 'Spring', 2: 'Summer', 3: 'Fall', 4: 'Winter'}
df['season_name'] = df['season'].map(season_map)

sns.boxplot(x='season_name', y='count', data=df, palette='viridis') # Using
# season names and a different palette
plt.title('Electric Cycle Demand Across Seasons', fontsize=16, weight='bold') #
# Increased title font size
plt.xlabel('Season', fontsize=14) # Changed xlabel
plt.ylabel('Rental Count', fontsize=14) # Increased ylabel font size
plt.grid(axis='y', linestyle='--', linewidth=0.5)
plt.tight_layout()
plt.show()
```

ANOVA Season: F-statistic = 236.95, P-value = 0.0000



## 27 Insights:

- The ANOVA test for demand across seasons resulted in a p-value of 0.0000. Because the p-value is much lower than 0.05, we reject the null hypothesis.
- This means there is a statistically significant difference in the average electric cycle demand across the different seasons.
- The graph shows that demand is highest in Fall (Season 3) and lowest in Spring (Season 1).

## 28 Recommendation (Season vs Demand):

- **Leverage Seasonality:** Since season significantly affects demand (as shown by the ANOVA p-value of 0.0000 and the boxplot), Yulu should adjust its operations based on seasonal patterns.
- Increase fleet size and marketing efforts in peak seasons (Fall and Summer) and consider promotions during off-peak seasons (Spring and Winter) to boost ridership.

## 29 Question 3: Is There a Significant Difference in Electric Cycle Demand Across Different Weather Conditions?

```
[50]: # Sort weather types
weather_types = sorted(df['weather'].unique())

# Create groups of 'count' for each weather type
```

```

weather_groups = [df[df['weather'] == w]['count'] for w in weather_types]

# Perform One-Way ANOVA
f_stat_weather, p_val_weather = f_oneway(*weather_groups)

# Print ANOVA result
print(f"ANOVA Test - Weather vs Demand:")
print(f"F-statistic = {f_stat_weather:.2f}")
print(f"P-value      = {p_val_weather:.4f}")

# Boxplot for visual analysis
plt.figure(figsize=(10, 6)) # Increased figure size
# Map weather numbers to names for better readability
weather_map = {1: 'Clear', 2: 'Mist/Cloudy', 3: 'Light Snow/Rain', 4: 'Heavy_
↳ Rain/Ice'}
df['weather_name'] = df['weather'].map(weather_map)

sns.boxplot(x='weather_name', y='count', data=df, palette='YlGnBu') # soothing_
↳ blue-green palette

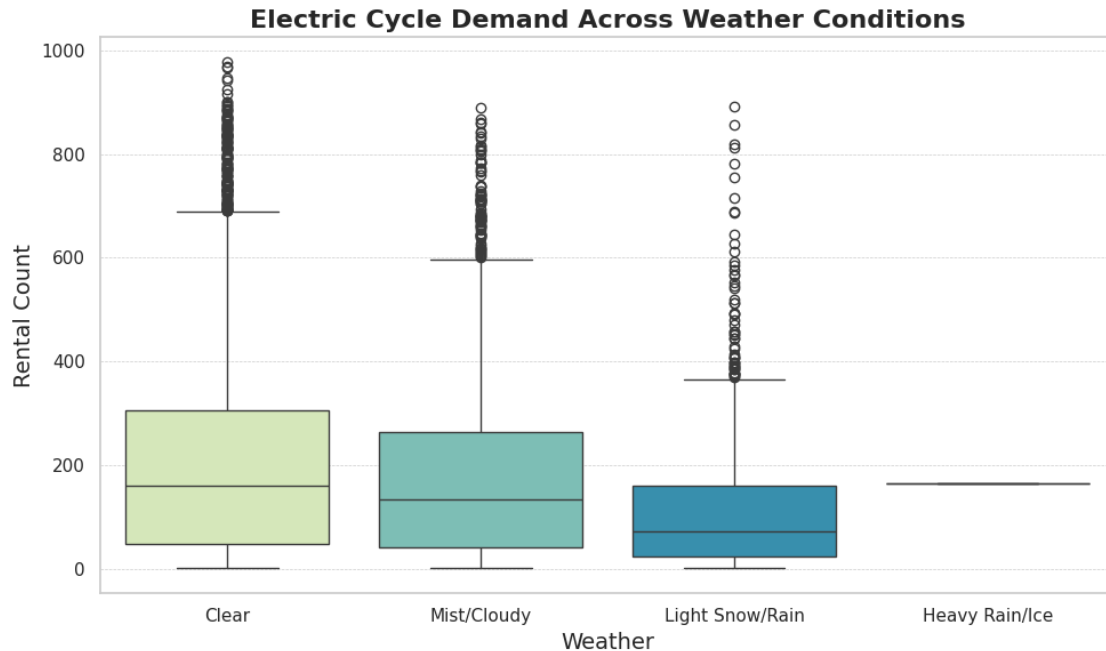
# Add labels and styling
plt.title('Electric Cycle Demand Across Weather Conditions', fontsize=16,
↳ weight='bold') # Increased title font size
plt.xlabel('Weather', fontsize=14) # Changed xlabel
plt.ylabel('Rental Count', fontsize=14) # Increased ylabel font size
plt.grid(axis='y', linestyle='--', linewidth=0.5)
plt.tight_layout()
plt.show()

```

```

ANOVA Test - Weather vs Demand:
F-statistic = 65.53
P-value      = 0.0000

```



### 30 Insights:

- The p-value from the test for weather vs. demand is 0.0000. Since the p-value is less than 0.05, we reject the null hypothesis, meaning weather significantly affects demand.
- Demand is highest in clear weather (Weather 1) and decreases as weather worsens, with very low demand in heavy rain/ice (Weather 4).

### 31 Recommendation:

- Weather depends on Season: Certain weather is more common in some seasons.
- Use seasonal weather patterns to predict demand better and plan resources accordingly.

### 32 Question 4: Is Weather Condition Statistically Dependent on Season?

```
[51]: # Create a contingency table
contingency_table = pd.crosstab(df['season'], df['weather'])

# Perform Chi-Square Test of Independence
chi2_stat, p_value, dof, expected_freq = chi2_contingency(contingency_table)

# Print results
print("Chi-Square Test - Season vs Weather")
print(f"Chi2 Statistic : {chi2_stat:.2f}")
```



```

print(f"P-value          : {p_value:.4f}")
print(f"Degrees of Freedom: {dof}")

# Visualize contingency table
plt.figure(figsize=(6, 4))
sns.heatmap(contingency_table, annot=True, fmt='d', cmap='YlGnBu', cbar=False)
plt.title("Season vs Weather - Frequency Heatmap", fontsize=14, weight='bold')
plt.xlabel("Weather Type", fontsize=12)
plt.ylabel("Season", fontsize=12)
plt.tight_layout()
plt.show()


```

Chi-Square Test - Season vs Weather

Chi2 Statistic : 49.16

P-value : 0.0000

Degrees of Freedom: 9



### 33 Insights:

- The Chi-Square test between season and weather gave a p-value of **0.0000**.
- Since the p-value is less than 0.05, we **reject the null hypothesis**. This means weather condition **depends on season**.
- The heatmap shows that certain weather types happen more in some seasons, supporting this link.


## 34 STEP 7: Assumption Checks (Optional)

```
[53]: # === Q-Q Plot to assess normality of 'count' variable ===
plt.figure(figsize=(6, 4))
scipy.stats.probplot(df['count'], dist="norm", plot=plt)
plt.title("Q-Q Plot for 'Count' (Demand)", fontsize=14, weight='bold')
plt.grid(True, linestyle='--', linewidth=0.5)
plt.tight_layout()
plt.show()

# === Levene's Test for Equal Variance Between Working and Non-Working Days ===
working = df[df['workingday'] == 1]['count']
non_working = df[df['workingday'] == 0]['count']

levene_stat, levene_p = levene(working, non_working)

# Print result
print("Levene's Test for Equal Variance Between Groups")
print(f"Levene Statistic : {levene_stat:.2f}")
print(f"P-value : {levene_p:.4f}")
```



```
Levene's Test for Equal Variance Between Groups
Levene Statistic : 0.00
P-value : 0.9438
```

## 35 Insights:

- The Q-Q plot shows that the 'count' data is not normally distributed.
- Levene's Test (P-value = 0.9438) indicates that the variances of 'count' are approximately equal between working and non-working days.

## 36 Recommendations

- Demand is not normal: Consider using non-parametric tests (like Mann-Whitney or Kruskal-Wallis) instead of or alongside t-test/ANOVA.
- Variances are equal (Working vs Non-Working): The standard t-test is appropriate for comparing these groups regarding variance assumptions.

```
[54]: print("""
Summary of Hypothesis Tests:
-----
Test          | Variable(s)          | P-Value | Inference
-----|-----|-----|-----
T-Test        | Workingday vs Count | {:.4f}   | {} workingday has effect
ANOVA         | Season vs Count     | {:.4f}   | {} season affects demand
ANOVA         | Weather vs Count    | {:.4f}   | {} weather affects demand
Chi-Square    | Season vs Weather   | {:.4f}   | {} weather depends on season
""").format(
    p_val,
    "Reject H0 =>" if p_val < 0.05 else "Fail to Reject H0 =>",
    p_val_season,
    "Reject H0 =>" if p_val_season < 0.05 else "Fail to Reject H0 =>",
    p_val_weather,
    "Reject H0 =>" if p_val_weather < 0.05 else "Fail to Reject H0 =>",
    chi2_p,
    "Reject H0 =>" if chi2_p < 0.05 else "Fail to Reject H0 =>"
))
```

Summary of Hypothesis Tests:

```
-----
Test          | Variable(s)          | P-Value | Inference
-----|-----|-----|-----
-
T-Test        | Workingday vs Count | 0.2264  | Fail to Reject H0 =>
workingday has effect
ANOVA         | Season vs Count     | 0.0000  | Reject H0 => season affects
demand
ANOVA         | Weather vs Count    | 0.0000  | Reject H0 => weather affects
demand
Chi-Square    | Season vs Weather   | 0.0000  | Reject H0 => weather depends
on season
```

## 37 Insights:

- The T-Test for **Workingday vs Count** has a P-value of **0.2264**. Since this is  $> 0.05$ , we **fail to reject the null hypothesis**. This means there is **no statistically significant difference** in the average electric cycle demand between working days and non-working days.
- The ANOVA for **Season vs Count** has a P-value of **0.0000**. Since this is  $< 0.05$ , we **reject the null hypothesis**. This means **season significantly affects** the average electric cycle demand.
- The ANOVA for **Weather vs Count** has a P-value of **0.0000**. Since this is  $< 0.05$ , we **reject the null hypothesis**. This means **weather significantly affects** the average electric cycle demand.
- The Chi-Square test for **Season vs Weather** has a P-value of **0.0000**. Since this is  $< 0.05$ , we **reject the null hypothesis**. This means **weather condition is statistically dependent on season**.

**Overall Conclusion:** Season and weather conditions are significant factors influencing electric cycle demand, while the distinction between working and non-working days does not show a statistically significant difference in average demand based on this analysis.

## 38 Recommendations

Based on the analysis, here are key recommendations for Yulu:

- **Leverage Seasonality:** Increase resources in peak seasons (Fall/Summer), consider promotions in off-peak (Spring/Winter).
- **Adapt to Weather:** Use dynamic pricing based on weather, efficiently redistribute bikes during poor conditions, and communicate with users.
- **Optimize for Commute:** Continue focusing on bike availability in commuter areas during peak hours.
- **Boost Non-Working Day Use:** Analyze weekend patterns and target promotions for leisure/recreational rides.
- **Predict Demand:** Consider building a model for more precise forecasts and resource planning.

This notebook was converted with [convert.ploomber.io](https://convert.ploomber.io)